

Assignment 3 - Homework Exercises on Streaming Algorithms

Duy Pham - 0980384
 Maciej Wikowski - 0927420
 Pattarawat Chormai - 0978675

October 22, 2015

Str.I-1

Suppose a deterministic ALG for ELEMENT UNIQUENESS uses at most s bits of storage. Then ALG can only be in 2^s states at any point in time, in particular after processing $m/2$ tokens.

On the other hand, the number of different frequency vector $F[1, \dots, n]$ where $F[j] = 1$ if an item j exists in the stream otherwise $F[j] = 0$ for all $j \in [n]$, that can be generated by a stream of $m/2$ tokens from $[n]$ is equal to the number of way we put $m/2$ balls into n bins :

$$\binom{n}{m/2} \geq \left(\frac{n}{m/2}\right)^{m/2} = 2^{m/2 \log(2n/m)}$$

Hence, when $s < (m/2) \log(2n/m)$, there must be two sequences $\sigma_1 := \langle a_1, \dots, a_{m/2} \rangle$ and $\sigma'_1 := \langle a'_1, \dots, a'_{m/2} \rangle$ whose frequency vectors are different but ALG is in the same state after processing σ_1 as it would be after processing σ'_1 . Now consider $\sigma_2 := \langle a_{m/2+1}, \dots, a_m \rangle$ which contains an item j in the stream. Such that all items in $\sigma_1 \circ \sigma_2$ are unique, while there is a duplicated item in $\sigma'_1 \circ \sigma_2$. Indeed such a stream exists if we take an item $j \in [n]$ for which $F_{\sigma_1}[j] = 0$ and $F_{\sigma'_1}[j] = 1$, and all items in σ_2 are distinct from σ_1 and σ'_1 except only j . Then j is unique in $\sigma_1 \circ \sigma_2$ but not in $\sigma'_1 \circ \sigma_2$.

Since ALG is deterministic and the state of ALG after processing σ_1 is the same as it would be after processing σ'_1 , we can conclude that ALG will report the same answer for the $\sigma_1 \circ \sigma_2$ and $\sigma'_1 \circ \sigma_2$. In fact, this is not correct, because j is unique in $\sigma_1 \circ \sigma_2$ but not in $\sigma'_1 \circ \sigma_2$. Hence, any deterministic streaming algorithm that solves ELEMENT UNIQUENESS exactly must use at least $\Omega(m \log(2n/m))$ bits of storage.

Str.I-3

The algorithm in vanilla model can be adapted to cash-register model by initialise $c(j)$, a counter of j , with c or increase $c(j)$ by c if $j \in I$. Secondly, if there are too many items in I , $|I| \geq 1/\epsilon$, $c(j)$ for all j in I will be decreased by $\min_{j \in I} c(j)$.

Next, let denote a token (j^*, c^*) as a token that is going to be processed next and j_{\min} whose $c(j_{\min}) = \min_{j \in I} c(j)$ in that time. We will argue that the algorithm in cash-register model correctly computes a superset of the ϵ -frequent items by comparing its behaviours with the original algorithm in vanilla model in 2 conditions, namely when $c^* \geq \minFreq$ and $c^* < \minFreq$.

For $c^* \geq \minFreq$, after the original algorithm processes a token $(j^*, 1)$ $c(j_{\min})$ time, j_{\min} will be removed from I and when it processes the rest of $(j^*, 1)$, $c(j^*)$ will equal to $c^* - c(j_{\min})$. This is exactly the same as the modified algorithm does when (j^*, c^*) arrives, j_{\min} will be removed since $c(j_{\min})$ becomes zero after subtracted by itself and j^* remains in I with $c(j^*) := c^* - c(j_{\min})$.

Algorithm 1 ϵ -Frequent Items

Require: A stream $\langle a_1, \dots, a_m \rangle$ in Cash Register Model where $a_i = (j, c)$

Initialize $I \leftarrow \emptyset$

Process(a_i):

if $j \in I$ **then**

$c(j) \leftarrow c(j) + c$

else

 Insert j into I with counter $c(j) = c$

if $|I| \geq 1/\epsilon$ **then**

$MinFreq \leftarrow \min_{j \in I} c(j)$

for all items $j \in I$ **do**

$c(j) \leftarrow c(j) - MinFreq$; delete j from I when $c(j) = 0$

end for

end if

end if

return I

For $c^* < minFreq$, after the original algorithm processes a token $(j^*, 1)$ c^* time, $c(j)$ for $j \in I$ will be subtracted by c^* and there is no j^* in I which is the same result from the modified algorithm does.

Therefore, we can conclude that the modified algorithm returns correct result.

Str.II-1

(i)

We know that $E[\tilde{\Phi}(\sigma)] = \Phi(\sigma)$ and $\text{Var}[\tilde{\Phi}(\sigma)] = (1/3) \cdot \Phi(\sigma)$, thus:

$$\Pr[|\tilde{\Phi}(\sigma) - \Phi(\sigma)| \geq c \cdot \Phi(\sigma)] = \Pr[|\tilde{\Phi}(\sigma) - E[\tilde{\Phi}(\sigma)]| \geq 3c \cdot \text{Var}[\tilde{\Phi}(\sigma)]]$$

According to Chebyshev Inequality, we have:

$$\begin{aligned} & \Pr[|\tilde{\Phi}(\sigma) - E[\tilde{\Phi}(\sigma)]| \geq 3c \cdot \text{Var}[\tilde{\Phi}(\sigma)]] \\ &= \Pr[|\tilde{\Phi}(\sigma) - E[\tilde{\Phi}(\sigma)]| \geq 3c \cdot \sqrt{\text{Var}[\tilde{\Phi}(\sigma)]} \cdot \sqrt{\text{Var}[\tilde{\Phi}(\sigma)]}] \\ &\leq \frac{1}{9 \cdot c^2 \cdot \text{Var}[\tilde{\Phi}(\sigma)]} \end{aligned}$$

To let $1/6$ be the upper bound of this probability, we have to choose a value of c such that:

$$\begin{aligned} 9 \cdot c^2 \cdot \text{Var}[\tilde{\Phi}(\sigma)] &= 6 \\ c &= \sqrt{\left(\frac{6}{9 \text{Var}[\tilde{\Phi}(\sigma)]}\right)} \\ &= \sqrt{\left(\frac{2}{3 \frac{\Phi(\sigma)}{3}}\right)} \end{aligned}$$

Since we know that $\Phi(\sigma) > 3$, then:

$$c = \sqrt{\left(\frac{2}{\Phi(\sigma)}\right)} < \sqrt{\left(\frac{2}{3}\right)}$$

This value of c satisfies the condition that $0 < c < 1$.

(ii)

The idea of the new algorithm is to run ALG k times, then taking the median of the k return values $\tilde{\Phi}(\sigma)$ (the Median trick). Before modifying the algorithm formally, we make some analysis as follows.

Let X_i be an indicator random variable, which is defined as:

$$X_i = \begin{cases} 1 & \text{if } \Pr[|\tilde{\Phi}(\sigma) - \Phi(\sigma)| \geq c \cdot \Phi(\sigma)] \\ 0 & \text{otherwise} \end{cases}$$

Let $X = \sum_{i=1}^k X_i$ be the random variable representing the final outcome of the new algorithm. We have:

$$\mathbb{E}[X] = \mathbb{E}\left[\sum_{i=1}^k X_i\right] = \sum_{i=1}^k \mathbb{E}[X_i]$$

In this algorithm, we used the specified value of c in part (i), which confirms that:

$$\Pr[|\tilde{\Phi}(\sigma) - \Phi(\sigma)| \geq c \cdot \Phi(\sigma)] \leq 1/6$$

which implies that

$$\Pr[X_i = 1] \leq 1/6$$

Because X_i is an indicator random variable, then:

$$\mathbb{E}[X_i] = 1/6$$

Thus, we have:

$$\mathbb{E}[X] = \sum_{i=1}^k \mathbb{E}[X_i] = k/6$$

When $X_i = 1$, it implies that we do not have the event described in the problem statement. The probability that we do not have it after performing the Median trick is:

$$\begin{aligned} \Pr[X > \frac{k}{2}] &= \Pr[X > 3 \mathbb{E}[X]] = \Pr[X > (1 + 2) \mathbb{E}[X]] \\ &\leq \left(\frac{e^2}{3^3}\right)^{\frac{k}{6}} \\ &= \left(\frac{e^2}{27}\right)^{\frac{k}{6}} \end{aligned}$$

Thus, the probability that we get the event described is $1 - \left(\frac{e^2}{27}\right)^{\frac{k}{6}}$. To get the desired value $(1 - \delta)$, we have to choose k such that:

$$\begin{aligned}
\delta &\geq \left(\frac{e^2}{27}\right)^{\frac{k}{6}} \\
\iff \frac{1}{\delta} &\leq \left(\frac{27}{e^2}\right)^{\frac{k}{6}} \\
\iff \log \frac{1}{\delta} &\leq \frac{k}{6} \log \frac{27}{e^2} \\
\iff k &\geq \frac{6}{\log \frac{27}{e^2}} \cdot \log \frac{1}{\delta}
\end{aligned}$$

We can choose

$$k = \lceil 4 \log \frac{1}{\delta} \rceil$$

In conclusion, we modify the algorithm as represented in algorithm 2.

Algorithm 2 Revised Algorithm Taking ALG as Sub-Routine

Require: A stream $\sigma = \langle a_1, \dots, a_m \rangle, \delta$

Ensure: Statistics $\tilde{\Phi}(\sigma)$

Operation:

Choose $k := \lceil 4 \log \frac{1}{\delta} \rceil$

Init $J := \emptyset$

for i from 1 to k **do**

$\tilde{\Phi}_i(\sigma) := \text{ALG}()$

$J = J \cup \tilde{\Phi}_i(\sigma)$

end for

return The median of set J

The algorithm needs to store the set J of k elements, and each run of ALG stores $B(n, m)$ bits, so the total number of bits used by the algorithm is $O(kB(n, m))$.

Str.III-1

In theorem 10.4, we prove the lower bound and the upper bound of the Count-Min sketch. To get the lower bound, we use the property of the Strict Turnstile model that $F_\sigma[j] \geq 0, \forall j \in [n]$, then $C[s, h_s(j)] \geq F_\sigma[j]$. In the general turnstile model, $F_\sigma[j]$ can be negative, so $C[s, h_s(j)]$ is not always greater than or equal to $F_\sigma[j]$. The lower bound is not $F_\sigma[j]$ anymore.

To get the upper bound, we use Markov Inequality. But it holds only if X is a non-negative random variable. In this case, $X = C[s, h_s(j)] - F_\sigma[j]$ can be negative, so the upper bound can also be different.

So in general, the proof of 10.4 does not work in the general turnstile model.

Str.III-2

According to the question we know that $\tilde{F}_\sigma[j] > m/2 + F_\sigma[j]$. That means $h(j)$ must be equal to $h(j^*)$ where j^* is an item that occurs $m/2 + 1$ time. We also know that the probability that $h(j) = h(j^*)$ is $1/k$. Thus $h(j) \neq h(j^*)$ is $(k-1)/k$.

Then the probability that the hash value of other $m/2 - 1$ items is not $h(j^*)$ is :

$$\begin{aligned}
\left(\frac{k-1}{k}\right)^{m/2-1} &< (1 - 0.99) \\
\left(\frac{9}{10}\right)^{m/2-1} &< 0.01 \\
(m/2 - 1) \log(0.9) &< \log(0.01) \\
m &> 2 \left(\frac{\log(0.01)}{\log(0.9)} + 1 \right) \\
m &> 89.41 \\
&\geq 90
\end{aligned}$$

Therefore, one possible stream for this particular case is a stream size 90 that has an item j^* occurs 46 time and one of the other 44 tokens is an item j whose $h(j) = h(j^*)$.